

CS100

Introduction to Programming

Tutorial 8: containers & threading

Part 1

std::unordered_map

Associative container

- Stores elements as as key-value pairs
- Similar to `std::map`, every element needs to have unique key

Search, insertion, and removal all have approximately **constant-time complexity!**

(remember that `map` `std::` has complexity of $O(\log(n))$)

```
name                employee
string              Employee

map<string, Employee *> employees;
```

Associative container

- Stores elements as key-value pairs
- Similar to `std::map`, every element needs to have unique key

Search, insertion, and removal all have approximately **constant-time complexity!**

(remember that `std::map` has complexity of $O(n)$)

name
`string`

employee
`Employee`

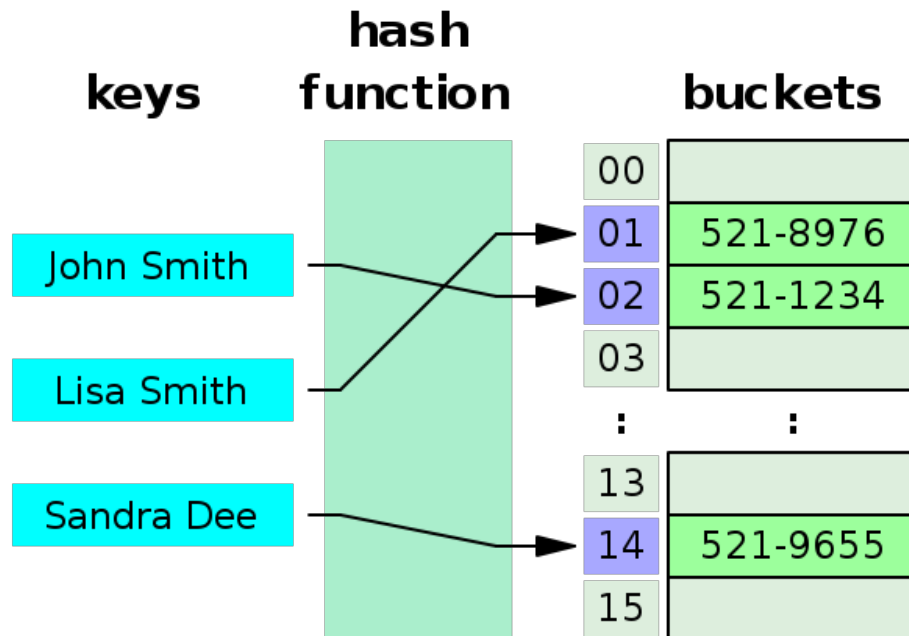
```
unordered_map<string, Employee *> employees;
```

How does it work?

- Internally, the elements are organized into buckets
- Which bucket an element is placed into depends on its key
 - From an arbitrary key-type, we derive a bucket-index
 - The bucket-index is called a hash
 - A hash-function is a function that maps an arbitrary input type to a defined type with defined range
- This allows fast access to individual elements, since once a hash is computed, it refers to the exact bucket the element is placed into

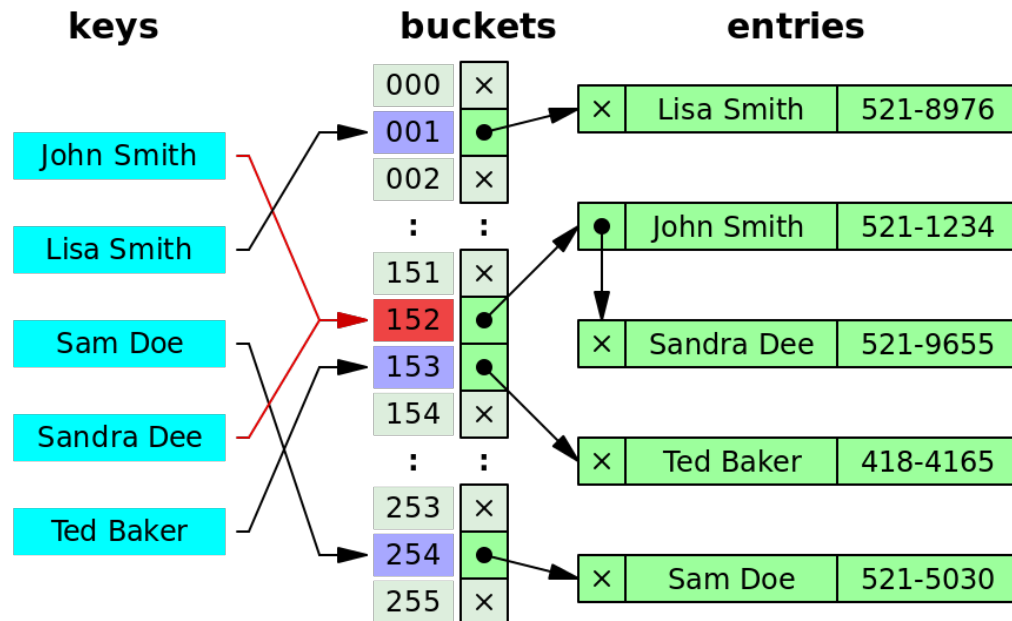
How does it work?

- Hash table
- https://en.wikipedia.org/wiki/Hash_table
- Ideally: Every key leads to an individual bucket



How does it work?

- Hash table
- https://en.wikipedia.org/wiki/Hash_table
- In practice: Limited number of bucket & extra collision resolving (e.g. chaining)



How does it work?

- Properties:

Algorithm	Average	Worst case
Space	$O(n)^{[1]}$	$O(n)$
Search	$O(1)$	$O(n)$
Insert	$O(1)$	$O(n)$
Delete	$O(1)$	$O(n)$

- How does the space property come across?
 - Only a table with (initially NULL) bucket pointers is installed
 - This table has a constant size equal to the range of the hash
 - The actual buckets are growing linearly with the actual number of elements in the list

Exercise

- Create a main function in which you
 - Fill a `std::map` with key value pairs of the form

`std::map<int, std::string>`
 - Notes:
 - The int can simply increase linearly for every element
 - The string can be same, dummy string everytime
- **Task 1:**
 - Add 10000 elements, then measure the time of adding 100 more elements

Exercise

- **Task 2:**

- Do the same for an unordered map
- What do you observe?

- **Task 3:**

- Now increase the initial size of the container by a factor of 10.
- What do you observe?

Part 2

MyList

Implement your own template list container

- Requirements:
 - Define a suitable type for the (doubly!) linked elements in the list
 - Implement a bidirectional STL-style iterator that iterates through the linked elements
 - Implement your own version of a template list
 - Required interface functions:

```
iterator begin();  
iterator end();  
bool empty();  
int size();  
void erase(iterator it);  
void insert(iterator it, const T& x);  
void push_back(const T& x);  
void clear();
```

Test MyList with the below code

- Instantiate new list and fill with elements

```
#include "MyList.hpp"
#include <iostream>
#include <algorithm>
#include <vector>

int main() {
    //instantiate one of my new list object and push_back two elements
    MyList<double> myList;
    myList.push_back(0.0);
    myList.push_back(1.0);

    ...
}
```

Test MyList with the below code

- Iterate through the elements and print them out

```
#include "MyList.hpp"
#include <iostream>
#include <algorithm>
#include <vector>

int main() {
    ...

    //use the iterator to loop through the elements
    //and print them in the console
    MyList<double>::iterator it = myList.begin();
    while( it != myList.end() ) {
        std::cout << *it << "\n";
        it++;
    }

    ...
}
```

Test MyList with the below code

- Copy the content over to an `std::vector` using `std::copy`!

```
#include "MyList.hpp"
#include <iostream>
#include <algorithm>
#include <vector>

int main() {
    ...

    //create a standard vector and copy all the elements over
    //(exploit the fact that our iterators are also STL iterators!)
    std::vector<double> vec;
    std::copy( myList.begin(), myList.end(), std::back_inserter(vec) );

    ...
}
```

Test MyList with the below code

- Verify that the content has been correctly copied!

```
#include "MyList.hpp"
#include <iostream>
#include <algorithm>
#include <vector>

int main() {
    ...

    //verify that the content has been correctly copied
    //over to the vector
    for( int i = 0; i < vec.size(); i++ )
        std::cout << vec[i] << "\n";
    return 0;
}
```


Part 3

**Implement two threads that
communicate with each other**

Synchronization between threads

- Apart from just protecting data, sometimes we may wish for one thread to wait until another thread has something done
- In C++:
 - Conditional variables
 - Futures

std::condition_variable

- A synchronization primitive that can be used to block a thread or multiple threads at the same time, until
 - A notification is received from another thread
 - A time-out expires

`std::condition_variable`

- A thread that intends to wait on `std::condition_variable` has to acquire a `std::unique_lock` first
- The wait operations atomically release the mutex and suspend the execution of the thread
- When the condition variable is notified, the thread is awakened, and the mutex is reacquired

Example

```
std::mutex mut;  
std::queue<data_chunk> data_queue;  
std::condition_variable data_cond;
```



Mutex to protect resource

```
void data_preparation_thread() {  
    while( more_data_to_prepare() ) {  
        data_chunk data = prepare_data();  
        std::lock_guard<std::mutex> lk(mut);  
        data_queue.push(data);  
        data_cond.notify_one();  
    }  
}
```

```
void data_processing_thread() {  
    while(true) {  
        std::unique_lock<std::mutex> lk(mut);  
        data_cond.wait(lk, []{return !data_queue.empty();});  
        data_chunk data = data_queue.front();  
        data_queue.pop();  
        lk.unlock();  
        process(data);  
        if(is_last_chunk(data))  
            break;  
    }  
}
```

Example

```
std::mutex mut;  
std::queue<data_chunk> data_queue;  
std::condition_variable data_cond;
```

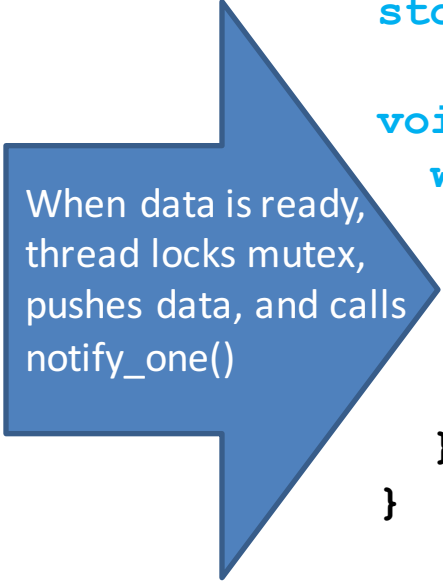


Queue used to pass data

```
void data_preparation_thread() {  
    while( more_data_to_prepare() ) {  
        data_chunk data = prepare_data();  
        std::lock_guard<std::mutex> lk(mut);  
        data_queue.push(data);  
        data_cond.notify_one();  
    }  
}
```

```
void data_processing_thread() {  
    while(true) {  
        std::unique_lock<std::mutex> lk(mut);  
        data_cond.wait(lk, []{return !data_queue.empty();});  
        data_chunk data = data_queue.front();  
        data_queue.pop();  
        lk.unlock();  
        process(data);  
        if(is_last_chunk(data))  
            break;  
    }  
}
```

Example



When data is ready,
thread locks mutex,
pushes data, and calls
notify_one()

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;

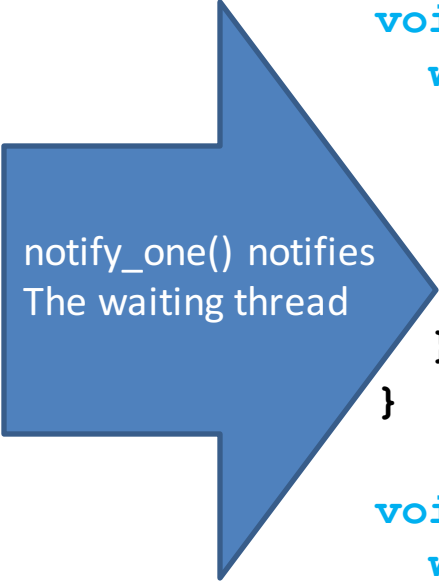
void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}

void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```

Example

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;

void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}
```



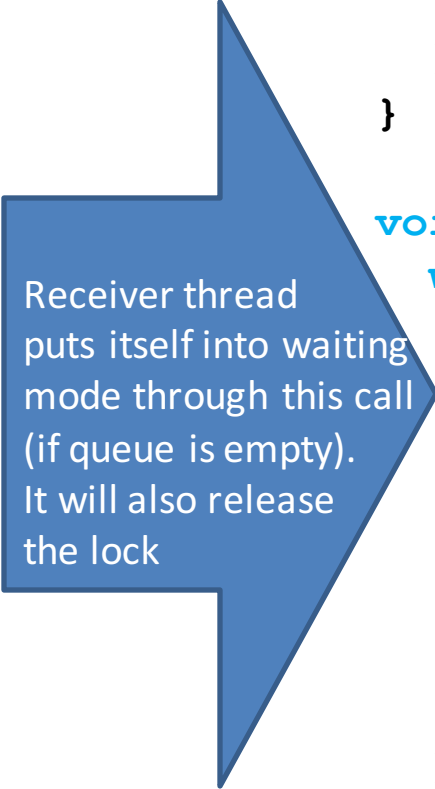
notify_one() notifies
The waiting thread

```
void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```


Example

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;
```

```
void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}
```



Receiver thread
puts itself into waiting
mode through this call
(if queue is empty).
It will also release
the lock

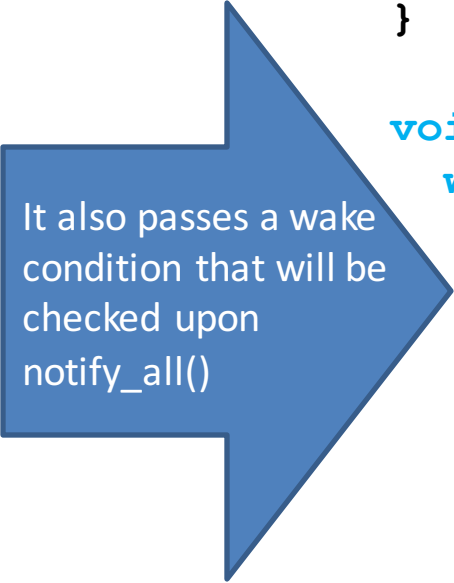
```
void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```

Example

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;
```

```
void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}
```

```
void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```



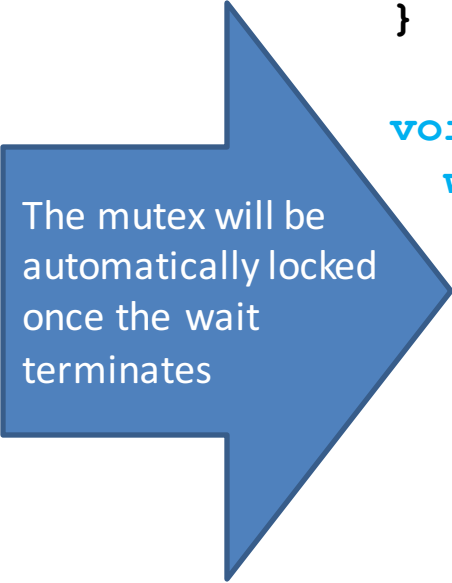
It also passes a wake condition that will be checked upon notify_all()

Example

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;
```

```
void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}
```

```
void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```



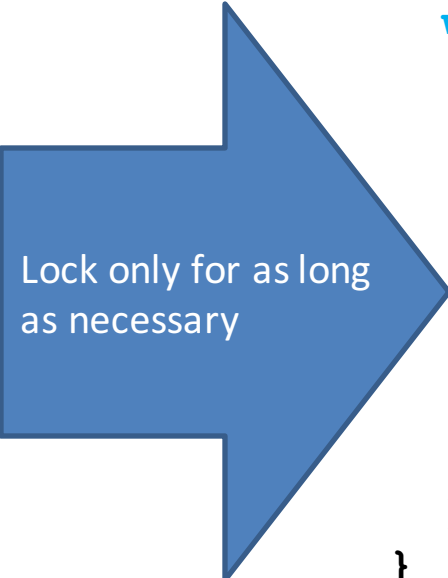
The mutex will be automatically locked once the wait terminates

Example

```
std::mutex mut;
std::queue<data_chunk> data_queue;
std::condition_variable data_cond;
```

```
void data_preparation_thread() {
    while( more_data_to_prepare() ) {
        data_chunk data = prepare_data();
        std::lock_guard<std::mutex> lk(mut);
        data_queue.push(data);
        data_cond.notify_one();
    }
}
```

```
void data_processing_thread() {
    while(true) {
        std::unique_lock<std::mutex> lk(mut);
        data_cond.wait(lk, []{return !data_queue.empty();});
        data_chunk data = data_queue.front();
        data_queue.pop();
        lk.unlock();
        process(data);
        if(is_last_chunk(data))
            break;
    }
}
```



Lock only for as long
as necessary

Playing with threads

- Implement two threads that are playing Ping-Pong
 - The first thread is configured to be the right-side player, and he says Ping if the ball is on the right side
 - The second thread is configured to be the left-side player, and he says Pong if the ball is on the left side
 - Realize the problem with a single function definition
 - Realize your implementation with a condition variable that sets the threads to sleep while they are waiting for the ball